
Trust and Safety of Human- and AI-Powered Visual Description: A Case Study of Be My Eyes

Natalie Grace Brigham and Tadayoshi Kohno

Be My Eyes provides a volunteer-powered visual description service that connects over one million blind and low vision users with more than ten million volunteers for everyday visual assistance tasks as well as Be My AI, an AI-powered counterpart. In this work, we qualitatively analyze 1,428 one-, two-, and three-star reviews from the Apple App Store and Google Play Store to holistically examine the concerns and harmful experiences of both users and volunteers. Our analysis surfaces online harassment and abuse (trolling, cyberflashing, reporting failures), privacy concerns beyond information exposure (data use for model training), and accessibility/safety problems (inaccurate descriptions, paternalistic guardrails). We additionally analyze platform policies to examine whether and how these issues are addressed by them. Based on challenges we find related to accessibility, security, privacy, and AI, we provide recommendations for strengthening trust and safety on visual description services that serve blind and low vision people's access needs.

1 Introduction

Be My Eyes offers a volunteer-powered visual description service (VDS) where blind and low vision people (BLV)¹ can request visual assistance from sighted volunteers through a mobile app on iOS and Android. Throughout this paper, we refer to these groups as “users” and “volunteers,” respectively. When a user requests assistance, the app notifies volunteers in an appropriate time zone who speak the user's language. Once connected, the user and available volunteer communicate on the app through a two-way audio feed and a one-way video feed from the user's device. Since 2023, the app has also offered Be My AI, a VDS powered by multimodal models.

Be My Eyes is the most prominent volunteer-powered VDS, widely included in academic

1. We recognize that preferences vary within disability communities regarding person-first and identity-first language (Hanson, Cavender, and Trewin 2015). In this paper, we use identity-first language, consistent with Sharif, McCall, and Bolante's (2022) recommendation.

literature (Alharbi, Brewer, and Schoenebeck 2022; Alharbi et al. 2024; Akter et al. 2020; Bennett et al. 2021; Gurari et al. 2018; Branham et al. 2017; Chang et al. 2025; Gurari et al. 2019; Gonzalez Penuela et al. 2024; Sharma et al. 2025; Stangl et al. 2020) and recommended in US government resources (Administration for Community Living 2025). The platform has received recognition including the 2021 Apple Design Award for Social Impact (Be My Eyes 2021a) and 2025 App Store Award for Cultural Impact (Be My Eyes 2025). As of April 2026, it reported having more than 10 million volunteers and more than a million users.

Previous work has investigated risks of sensitive information exposure through VDS use, such as unintentional disclosure of personally identifiable information in image or video backgrounds, and how users navigate these risks, such as choosing specific platform types for particular tasks (Gurari et al. 2019; Akter et al. 2020; Stangl et al. 2020; Alharbi, Brewer, and Schoenebeck 2022). We extend this work by stepping back and holistically considering trust and safety on Be My Eyes, including risks from user-volunteer interactions, platform data practices, and AI integration. Specifically, we examine:

RQ1: What are the trust and safety concerns of users and volunteers using Be My Eyes, and what harmful experiences have users and volunteers encountered?

To answer this question holistically across the platform's history, we qualitatively analyzed all 1,428 one-, two-, and three-star English-language reviews of Be My Eyes from the Apple App Store and Google Play Store, spanning 2015 through 2025. We believe reviews with these ratings sufficiently surface relevant concerns and experiences to answer RQ1; we provide an in-depth discussion of this scope in Section 3.1.

Our analysis surfaces trust and safety issues related to online harassment and abuse (e.g., trolling, cyberflashing), privacy concerns beyond information exposure detailed in prior work (Akter et al. 2020; Alharbi, Brewer, and Schoenebeck 2022; Stangl et al. 2020; Gurari et al. 2019) (e.g., data use for model training), and accessibility/safety (e.g., inaccurate descriptions). These issues span Be My Eyes in general, as well as apply specifically to the volunteer-powered service and Be My AI.

Mitigation of (potential or actual) harms to users and volunteers, or their concerns, can be executed in practice through design and implementation, or policy. Given that users and volunteers have consistently reported concerns and experiences related to these issues across multiple years, including recently, such concerns appear inadequately addressed through practical measures alone. Accordingly, we examine how the platform addresses—or fails to address—these issues at the policy level:

RQ2: What is the relationship between concerns and experiences surfaced in RQ1 and what is stated in Be My Eyes' privacy policies and terms of service?

To answer this question, we analyze three versions of Be My Eyes' terms of service and

four versions of its privacy policy, spanning from before Be My AI's March 2023 release through April 2026. Our analysis reveals varying levels of policy response to identified risks, as well as new concerns for users and volunteers introduced through data collection and sharing practices.

Based on tensions identified when answering RQ1 and RQ2 between (1) users' and volunteers' experiences and concerns, and (2) the service and its policies, we ask:

RQ3: What are fundamental trust and safety challenges between users' and volunteers' needs and the apparent functionality, performance, and business goals of Be My Eyes?

We identify challenges related to developing robust, non-obstructive, non-paternalistic safeguards and trauma-informed support as well as balancing the benefits of data collection and sharing with privacy risks, especially in this application's context. Based on these challenges, we provide guidance and recommendations for both Be My Eyes and VDS platforms in general.

Our work is the first to combine analysis of app reviews with terms of service and privacy policies to holistically examine trust and safety on a VDS platform. While grounded in Be My Eyes, the challenges we identify—related to online harassment and abuse, privacy violations, AI hallucinations and paternalistic content restrictions, and extractive data practices—reflect structural issues likely present across VDS platforms. Through our findings and recommendations, we aim to inform the development of VDSs that meet BLV people's access needs while appropriately leveraging volunteer contributions and protecting both groups from harm.

2 Background & Related Work

In this section, we provide background information about the company and app Be My Eyes (Section 2.1) and review prior work related to VDSs (Section 2.2).

2.1 Be My Eyes

The Be My Eyes mobile app was launched on iOS in 2015 to connect BLV people with volunteers who could provide visual assistance for everyday tasks (Be My Eyes, [n.d.-a](#)). Be My Eyes later introduced an Android app in 2017 (Be My Eyes [2021b](#)). As of November 2025, the Be My Eyes app had 11k reviews with an average rating of 4.8 on the Apple App Store and 32.3k reviews with an average rating of 4.5 on the Google Play Store.

In March 2023, Be My Eyes announced a partnership with OpenAI to introduce Be My AI, a service offering image descriptions powered by multimodal models (Be My Eyes [2023b](#)). Be My Eyes continued its investment in AI technologies with an October




Figure 1: Screenshots of Be My Eyes’ interface from the app’s Google Play Store page (in August 2026).

2024 partnership with Microsoft, providing video data from user-volunteer interactions to “train more inclusive models” (Be My Eyes 2024b). In another collaboration, they integrated with Ray-Ban Meta Smart Glasses to provide hands-free visual description calls in September 2024 (Be My Eyes 2024a).

Understanding users’ and volunteers’ onboarding and everyday experience on the app provides context for interpreting their concerns and experiences (RQ1). We provide that context here. Users and volunteers follow a similar onboarding process after installing the app:

1. Select role: “I need visual assistance” or “I’d like to volunteer” (see Figure 1a)
2. Agree to terms of service and privacy policy, which can (but are not required to) be read on the web by selecting the links provided within the app
3. Create account with first and last name, email, and password
4. Verify email by clicking link in verification email
5. Select preferred language
6. Grant permissions: notifications, microphone, camera (for users)

On the app’s home page, volunteers can optionally complete a short tutorial about answering calls and providing visual assistance. For users, a large, prominent “Call a volunteer” button initiates a video call (see Figure 1b). Through the menu bar, users can access the Be My AI feature to receive AI-generated visual descriptions of images (see




Figure 2: Flow chart of the processes for a user-volunteer call.

Figure 1c) and the “Service Directory” to connect with specific companies or service providers. Additionally, both users and volunteers can browse the “Community” tab, which contains advertisements, blog posts, news articles, and podcasts. Users may use assistive technologies such as magnifying software and screen readers to navigate and interact with the app.

Figure 2 shows the process for a user-volunteer call: the user selects “Call a volunteer,” volunteers are notified, and the first volunteer to answer is connected. During the call, there is a one-way video stream from the user’s back-facing camera and a two-way audio stream. Either party can end the call and after it concludes, both the user and volunteer can optionally provide call ratings or report issues. Both parties can reach support through these post-call mechanisms or via the company’s support email.

To answer RQ3, we make inferences about Be My Eyes’ goals, grounded in the company’s self-description and policies: The company is “committed to providing [their] tools and services to the blind and low-vision community without charging them” (Be My Eyes, n.d.-a) and, consequently, seeks alternative ways to sustain its business without charging users. Given their “mission is to make the world more accessible for people who are blind or have low vision... [Be My Eyes] incorporate[d] AI into the existing Be My Eyes platform” (Be My Eyes 2023b). Furthermore, the company contributes data to AI model training “to improve the accuracy and precision of scene understanding and descriptions, with the goal of increasing the utility of AI for the blind and low-vision community” (Be My Eyes 2024b).

2.2 Visual Description Services (VDSs)

Visual description services (VDSs) are used by BLV people for a variety of tasks, including reading physical text or writing, identifying items, describing qualities like color or state

(Brady et al. 2013), finding things, shopping assistance (Avila et al. 2016), fashion advice (Burton et al. 2012), and personal safety (Branham et al. 2017). These services can be powered by people, either paid agents (Aira, n.d.) or volunteers (Be My Eyes, n.d.-b), or by AI models (Microsoft, n.d.-a), and users' choice of platform is influenced by personal and contextual factors (Akter et al. 2020; Stangl et al. 2020).

Early academic work in VDS development includes VizWiz, developed by Bigham et al. (2010), which leveraged crowdworkers to answer visual questions. This work spawned substantial follow-on research, including the development of a dataset containing 31,000 visual question-answer pairs that supports AI development for automated visual question answering (Gurari et al. 2018). A subset of 5,537 images from this dataset, annotated with regions showing private content, was also curated to better understand VDS users' privacy disclosures and to encourage the development of algorithms that can assist in preventing their unintended disclosures (Gurari et al. 2019).

Trust and safety with VDSs. Prior work reveals nuanced and context-dependent privacy considerations in VDS use. In a survey of 155 BLV participants, Akter et al. (2020) found that privacy concerns varied based on the type of assistance (AI, volunteer, or paid agent) and the information involved in each task. Participants expressed particular concern about inadvertently exposing bystanders' faces and revealing negative self-presentation, such as messy spaces. Similarly, Stangl et al. (2020) examined attitudes toward sharing different types of private visual content (e.g., financial information, medical information, names, locations) and found that these attitudes inform users' choice of VDS types (e.g., AI- or volunteer-powered).

Given concerns about unintended information disclosure, researchers have investigated visual obfuscation techniques to mitigate privacy risks. However, through interviews with VDS users, Alharbi, Brewer, and Schoenebeck (2022) found that existing obfuscation approaches do not adequately address the needs of BLV people. Building on these findings, Zhang et al. (2024) prototyped and co-designed a privacy-preserving visual assistance system. Alharbi et al. (2025) evaluated AI-enabled visual obfuscation technology, finding that textual descriptors of what is obfuscated were insufficient to support blind people's error detection, and developed guidance for improved descriptors and the co-design of such systems.

Additionally, through interviews with 10 blind participants, Branham et al. (2017) found that participants were interested in using VDSs to enhance their physical safety, for example by identifying potential threats or people with weapons in their vicinity. Ahmed et al. (2015) and Ahmed et al. (2016) identified privacy and security concerns of BLV people in both online and offline contexts, recommending that future assistive technologies (e.g., VDSs) include the ability to count and recognize people around them as well as leverage wearable hardware.

We build on this body of work by examining trust and safety concerns while using a VDS

beyond information disclosure and leakage, including online safety violations and data privacy issues, that require mitigations beyond obfuscation.

AI for visual description. AI technologies have been used in visual description well before the 2023 surge in generative AI popularity and availability, particularly for text-reading tasks. Microsoft released Seeing AI in 2017 (Microsoft, [n.d.-b](#)), which became a widely used VDS for reading visual text (Akter et al. [2020](#); Stangl et al. [2020](#); Alharbi, Brewer, and Schoenebeck [2022](#)).

In recent years, the advancements in multimodal models have enabled broader visual description capabilities. Several studies have examined BLV people's experiences with and attitudes toward these AI tools for visual description beyond text-based scenarios. Gonzalez Penuela et al. ([2024](#)) conducted a two-week diary study with 16 participants using an AI scene description tool, finding that the generated descriptions needed improvement to provide a satisfying and trustworthy experience. They later conducted a similar study with 20 users of multimodal large language model-enabled visual interpretation tools, who reported that the tools were trustworthy and somewhat satisfying, although they often produced incorrect answers or abstained from answering (Gonzalez Penuela et al. [2026](#)). Alharbi et al. ([2024](#)) found that BLV people employ multiple strategies to verify AI outputs and contest errors: leveraging non-visual skills, strategically consulting sighted people, and cross-referencing information across different devices and systems. Their work identifies shortcomings in AI VDSs current at the time and highlights opportunities for developing more accessible, explainable AI systems. Chang et al. ([2025](#)) found that current live video AI technologies effectively provide answers to BLV users for static visual scenes but have shortcomings in delivering essential live descriptions required in dynamic situations. Xie et al. ([2025](#)) conducted interviews with 14 users of Be My AI and found the system is limited by hallucinations and misinterpretations of social contexts, styles, and human identities as well as failure to grasp users' intentions.

To address these limitations, researchers have proposed design recommendations and new systems for more trustworthy, controllable visual description. Through interviews with BLV people, Sharma et al. ([2025](#)) provided design recommendations for generative AI tools, with particular focus on privacy and information handling. Xie et al. ([2025](#)) developed recommendations for systems that connect AI- and volunteer-powered VDS. Zamiri Zeraati et al. ([2026](#)) explored how blind users control responses in AI visual question answering systems, finding the need to support on-demand verbosity control, user-centered spatial references, proactive camera guidance, and verification mechanisms. WorldScribe (Chang, Liu, and Guo [2024](#)) is an automated live visual description service that tailors descriptions to users' intents and context and is adaptive to visual and sound contexts, and TouchScribe (Chang et al. [2026](#)) is a complementary system that provides real-time, adaptive feedback for hand-object interactions.

3 Methods

First, as described in Section 3.1, we collected and qualitatively analyzed all 1,428 one-, two-, and three-star English-language reviews for Be My Eyes from the Apple App Store and Google Play Store posted from the app's launch in 2015 and 2017, respectively, through November 2025. The app serves both users and volunteers through the same platform, though with different functionality for each group, so reviews from users and volunteers appear together. Then, as described in Section 3.2, we collected and analyzed multiple versions of Be My Eyes' terms of service (TOS) and privacy policy. For the policy documents, we used three versions of the TOS (from September 2017 through April 2026) and four versions of the privacy policy (from May 2020 through April 2026). These time frames capture the platform before and after the March 2023 announcement of Be My AI, through to current policies at our time of writing (April 2026). In Section 3.3, we discuss our positionality in conducting this research. As discussed in Section 4, our IRB determined this work does not constitute human subjects research.

3.1 App Reviews

One author conducted a preliminary analysis by reading reviews on the Apple App Store and Google Play Store for the Be My Eyes app, determining that concerns and harmful experiences appeared to surface primarily through one-, two-, and three-star (out of five stars) reviews. During analysis (described below), all codes were identified in the one-star reviews; the two- and three-star reviews did not introduce any new codes. Thus, while it is possible some concerns may not be captured in our findings—and we acknowledge this as a limitation—our analysis surfaces broad categories of issues not previously addressed in the literature, with one-star reviews providing the core codebook and two- and three-star reviews offering additional depth and validation.

Furthermore, in the Apple App Store, one-, two-, and three-star reviews made up about 4.1% of reviews and 3.0% in the Google Play Store. The large number of users and volunteers (over 11 million combined) means both that a vast majority of people have a primarily positive experience on the app and that small percentages of negative experiences constitute large absolute numbers (Robertson 2022). Additionally, online abuse can go unreported (Kass et al. 2026; Duggan 2017), meaning actual rates might be higher.

We chose to capture these reviews from the app's history because it provides the most comprehensive view of user and volunteer concerns and experiences over time and insight into the persistence of certain issues. However, for readers interested only in the current state and functionality of the Be My Eyes app, our decision to analyze all reviews over time may be a limitation.

Table 1: Collected review counts.

Star rating	Google Play Store		Apple App Store		Total	
	n	%	n	%	n	%
1	475	48.8	236	52.0	711	49.8
2	205	21.0	100	22.0	305	21.4
3	294	30.2	118	26.0	412	28.9
1–3	974		454		1,428	

Data collection. We manually collected all one-, two-, and three-star English-language reviews from the Apple App Store and Google Play Store in November 2025. The resulting review counts are reported in Table 1. The reviews are identified by IDs A1–A454 and P1–P974 for the Apple App Store and Google Play Store, respectively.

Qualitative analysis. To conduct our qualitative analysis, one researcher read through all of the reviews and developed a preliminary codebook based on concerns and experiences expressed in the reviews. Our approach was inductive as we sought to center the perspectives of volunteers and users, as discussed in Section 3.3. Codes were non-exclusive, allowing reviews to receive multiple codes when applicable.

The researcher and an external researcher then independently applied the codebook to a random sample of 10% of the collected one-, two-, and three-star reviews from each store (Apple App Store = 45; Google Play Store = 97). The researchers had only minor discrepancies (91% agreement), which involved ambiguous, few-word reviews and reviews that mentioned language support issues. To resolve these discrepancies, ambiguous reviews that could be interpreted in multiple, diverging ways were coded as “Unintelligible/Unclear concern,” and a new “Language” code was added for language support issues. After this clarification and modification of the codebook, both researchers’ codes were fully aligned. The first researcher then coded the remaining reviews for consistency. The final codebook with counts is available in Table 2.

3.2 TOS & Privacy Policies

To address RQ2 of how Be My Eyes’ policies related to the concerns and experiences surfaced in the app reviews, we analyzed Be My Eyes’ TOS and privacy policies.

Data collection. In December 2025 and again in February 2026, we collected the TOS and privacy policy from Be My Eyes’ website. Using the Internet Archive’s Wayback Machine, we collected earlier versions of the TOS and privacy policy to ensure we had at least one version of each from before the announcement of Be My AI in March 2023. In total, we collected privacy policies that went into effect on May 18, 2020² (PP1),

2. <https://web.archive.org/web/20230701120214/https://bemyeyes.com/privacy>

Table 2: Review codebook with frequencies for Google Play Store and Apple App Store reviews.

Code	Description	Google		Apple		Total	
		n	%	n	%	n	%
Technical	Errors, failures, malfunctions, missing features, or other technical problems	552	56.7	199	43.8	751	52.6
Call frequency	Volunteer reports insufficient call volume or infrequent matching with users	173	17.8	151	33.3	324	22.7
Unintelligible/ Unclear concern	Review contains concerns that cannot be fully understood or interpreted, or does not indicate a reason for negative rating	109	11.2	31	6.8	140	9.8
Support	Problems encountered when attempting to send feedback to, contact, or receive help from user support	56	5.7	8	1.8	64	4.5
Cyberflashing/ Sexual harassment	User displays or shares sexually explicit content on video calls with volunteers without consent	13	1.3	17	3.7	30	2.1
Other	Review contains a clear, specific concern that does not fit into any existing category (e.g., incorrect assumptions about pricing)	22	2.3	6	1.3	28	2.0
TOS & Privacy Policy	Issues or concerns related to Be My Eyes' policies	22	2.3	5	1.1	27	1.9
Suspected trolling	Volunteer or user appears to be engaging in disruptive behavior for amusement or to cause upset	8	0.8	15	3.3	23	1.6
Hostile/ Uncomfortable interaction	Interaction between user and volunteer is unpleasant or contentious, but does not meet the threshold for trolling or fraud	6	0.6	17	3.7	23	1.6
Accessibility	Problems related to app accessibility features or barriers	18	1.8	4	0.9	22	1.5
AI-related	Issues related to Be My AI	14	1.4	7	1.5	21	1.5
Language	App does not support user's language or has translation/localization problems	13	1.3	3	0.7	16	1.1
Account suspension	Complaints or inquiries related to suspended or restricted accounts	9	0.9	6	1.3	15	1.1
Misunderstanding of visual impairment	User misunderstands the app's purpose or target audience, such as questioning how blind people can use technology	9	0.9	5	1.1	14	1.0
Marketing material	Complaints about excessive marketing communications or notifications	9	0.9	3	0.7	12	0.8
Regional availability	App is not available in user's geographic region or country-specific restrictions prevent access or participation	7	0.7	1	0.2	8	0.6
Suspected fraud	User appears to be attempting to gain unauthorized information or access through deceptive means	6	0.6	2	0.4	8	0.6
Volunteer availability	User experiences long wait times or difficulty connecting with an available volunteer	3	0.3	4	0.9	7	0.5

September 18, 2023³ (PP2), July 25, 2024⁴ (PP3), and February 12, 2026⁵ (PP4), and TOS that went into effect on September 12, 2017⁶ (TS1), September 18, 2023⁷ (TS2), and February 12, 2026⁸ (TS3). TS3 and PP4 are the most recent versions of these documents at the time of writing (April 2026).

Qualitative analysis. We developed an a priori codebook based on the concerns and experiences identified in RQ1, following methodologies from prior work that analyzed parental control (Feal et al. 2020) and electronic monitoring (Owens et al. 2022) apps' privacy policies. One researcher applied this codebook to all collected policy documents. The codes were:

- What are the stated expectations and prohibited behaviors for *volunteers*?
- What are the stated expectations and prohibited behaviors for *users*?
- What are the data-sharing practices with third parties?
- Is video data used for AI training?
- Is image data used for AI training?

3.3 Positionality Statement

Drawing on feminist theory, we recognize the important role of reflexivity and positionality in our work (Bardzell and Bardzell 2011; Liang, Munson, and Kientz 2021; Darwin Holmes 2020). Our team consists of sighted researchers specializing in computer security and privacy, human-computer interaction, and AI ethics. As we research the experiences of users and volunteers on Be My Eyes, we acknowledge that we are outside of the BLV community and our interpretations and understanding of users' experiences are necessarily limited by this positionality. While not a substitute for community-based perspectives, we had an external researcher specializing in accessibility and with experience working with the BLV community conduct a sensitivity read.

We take inspiration from Garland-Thomson's (2002) feminist disability studies framework and Hamraie and Fritsch's (2019) Crip Technoscience commitments, particularly their emphasis on centering disabled people as designers and decision-makers in technology development. By building on the experiences and concerns of users and volunteers, our work seeks to provide insights and recommendations that honor their perspectives and support community-led development.

3. <https://web.archive.org/web/20240620183249/https://bemyeyes.com/privacy>

4. <https://web.archive.org/web/20251209035007/https://bemyeyes.com/privacy-policy/>

5. <https://web.archive.org/web/20260213013304/https://bemyeyes.com/privacy-policy/>

6. <https://web.archive.org/web/20230701120130/https://bemyeyes.com/terms>

7. <https://web.archive.org/web/20250406124903/https://bemyeyes.com/terms-of-service/>

8. <https://web.archive.org/web/20260317232423/https://bemyeyes.com/terms-of-service/>

4 Ethical Considerations

Our goal is not to discourage use of Be My Eyes, which provides free and valuable accessibility services to over a million users. Rather, we seek to surface trust and safety challenges at the intersection of accessibility, security, privacy, and AI and provide guidance and recommendations for navigating them. We hope these findings contribute to the strengthening of trust and safety on VDS platforms broadly.

Additionally, regarding our collection and use of public review data, we received an IRB determination from our institution that this research does not constitute human subjects research under federal regulations. Nevertheless, we considered Buck and Ralston's (2021) ethical guidelines for using social media data, focusing our analysis on patterns of concerns and experiences rather than on identifying or characterizing individual users. We include exact review text for two reasons: (1) reviewers likely expect their public reviews to be used by others evaluating the app, and (2) we aim to honor BLV reviewers' descriptions of their experiences without reinterpreting them through our positionality. Furthermore, we do not provide usernames for review authors in our results, given that they may contain names or other identifying information.

We extend these considerations to other publicly available content in our study, specifically a TikTok video from Be My Eyes' account. We link directly to this video where we discuss it (Section 5.2) since it is promotional content intended for the public.

5 Results

In this section, we explore users' and volunteers' concerns and harmful experiences (RQ1) and platform policies' connections to them (RQ2), organized in three groups. First, we cover online harassment and abuse (Section 5.1) and privacy (Section 5.2). Then, we examine challenges related to the Be My AI feature (Section 5.3). At the end of each subsection, we raise key challenges based on our findings (RQ3).

We provide excerpts from our data identified by the IDs detailed in Section 3. Apple App Store reviews are reported with the "yr(s) ago" label displayed when they were collected in November 2025 and Google Play Store reviews are reported with their precise date. We provide a label of "User" or "Volunteer" when the author's role is evident from the content of their review.

5.1 Online Harassment & Abuse

Reviews revealed online harassment and abuse on Be My Eyes spanning trolling, fraud, and sexual harassment, including cyberflashing. Consistent with patterns in broader online contexts (Lenhart et al. 2016), these behaviors manifested in various forms from different perpetrators.

Trolling & fraud. Volunteers experienced what Lenhart et al. (2016) term “direct harassment,” behaviors “that people do directly to one another,” including suspected trolling:

... I have received only one call and it was prankster pretending to be blind. They started swearing and verbally abusing me while I was with my children. No way to report them either. Disappointed that people are using this as a joke ... [P498, Volunteer, 11/21/2021]

Very upset with my first call. It's sad to see people are downloading this and using it to play around. It's sickening and just plain wrong! I downloaded the app 2 days ago and got my first call from a person who thought it was funny to to joke around ... [A149, Volunteer, 5 yrs ago]

and suspected fraud:

... recently SCAMMERS are using this app. And asking for WhatsApp numbers ... [P148, Volunteer, 12/27/2023]

Trolling and fraud are relatively common on large online platforms (Lenhart et al. 2016; Gottfried, Park, and Anderson 2025; Houtti et al. 2024). However, when users encounter these behaviors, they might experience compounded harms due to their disability: First, these attacks could constitute ableist harassment, whether intentional or not, by targeting disabled people seeking access. Second, they could function as “denial of access,” which occurs when a malicious actor uses platform features to prevent someone from accessing essential digital tools (Lenhart et al. 2016). On Be My Eyes, trolling volunteers effectively deny users access to the visual assistance the platform provides. For example, users and their families describe what they suspect was deliberately false information:

... my grandpa has been using this app, but when he calls people to help him a lot of the times, people give him the wrong information. we've had a lot of hecklers call him to tell him things that aren't true, and make his life even harder ... [A147, (Family member of) User, 5 yrs ago]

While some volunteers may provide incorrect information unintentionally, the presence of any trolling undermines the reliability of the service, creating additional barriers that prevent meaningful access.

Sexual harassment & cyberflashing. Sexual harassment was the most commonly reported form of online abuse in reviews. Volunteers primarily reported experiencing cyberflashing, nonconsensual digital exposure to sexual images or acts (McGlynn and Johnson 2021):

Answered a call and man inappropriately started pleasuring himself. Deleted immediately. [P216, Volunteer, 6/30/2023]

All reports of cyberflashing in our dataset appeared to come from volunteers, likely because only users can initiate video calls on the platform, enabling adversaries posing as users to expose themselves via camera. However, users may still face sexual harassment through verbal threats or attacks, such as comments about their physical appearance visible to volunteers during calls.

Platform non-liability. Although Be My Eyes' terms of service (TOS) explicitly prohibits many harmful behaviors, including trolling, fraud, and sexual harassment, it disclaims liability for harm resulting from user-volunteer interactions, stating:

Be My Eyes provides the platform to power and support these App Services. We do not create or pre-screen the videos, images, or other data transmitted through our App Services. We are not and cannot be responsible for the content being shared using our App Services. [TS3]

Pre-screening millions of users and volunteers is both logistically difficult and unlikely to prevent all abuse, and additional vetting may create barriers to providing the accessibility service the app seeks to provide. Nonetheless, our findings reveal substantial limitations in the current approach.

Platform response & reporting failures. The platform has reporting systems in place for violations of its TOS. However, online harassment and abuse are compounded by failures in these reporting systems:

... I received two harassment calls at night, both trying to expose their genitals to me and I had trouble reporting them after immediately hanging up. The app stopped responding when I tried to report those users ... [P10, Volunteer, 9/15/2025]

Issues with reporting appeared in the Google Play Store as early as January 2018:

... I wanted to send this via "Send Feedback," but it's programmed to Share instead ... [P951, Volunteer, 1/5/2018]

We identified 33 Google Play Store reviews, spanning from January 2018 [P951] to September 2025 [P5], stating that negative feedback and reporting features did not work on Android devices. We did not find comparable reviews in the Apple App Store.

The persistence of these reporting failures over seven years, combined with the range of harmful experiences documented in our findings, demonstrates a critical gap in platform support. Beyond functional reporting mechanisms, the potentially traumatic nature of experiences we identified, including cyberflashing, targeted harassment, and denial of access, necessitates comprehensive, trauma-informed support systems.

Trust and safety challenges. Trolling, fraud, and cyberflashing are distinct behaviors, but they share a common consequence of keeping volunteers and users from delivering and receiving, respectively, the assistance the service exists to provide. Addressing this pattern raises a first, proactive challenge:

Challenge 1: *Robust, non-obstructive safeguards*

Designing safeguards against trolling, fraud, and sexual harassment that do not inadvertently create additional access barriers for users or participation barriers for volunteers.

Yet safeguards alone cannot eliminate abuse, and our findings show that when it does occur, users and volunteers are often left without recourse as Be My Eyes disclaims liability for user-volunteer interactions [TS3], and 33 reviews spanning 2018–2025 describe reporting tools that do not work on Android devices. The combination of documented abusive behaviors, non-liability, and broken reporting mechanisms makes reactive support just as urgent as prevention, raising the challenge of:

Challenge 2: *Trauma-informed reporting and support*

Developing trauma-informed reporting and support systems that address the full spectrum of harms users and volunteers face, recognizing how denial of access for users may compound the impact of harassment and abuse.

5.2 Broadening Privacy Considerations

Prior work on VDS users' privacy concerns finds that users worry about unintentionally sharing conventionally sensitive information (e.g., medical or financial details), information affecting self-presentation (e.g., visible clutter), and information about others (e.g., bystanders' faces). Users also worry about necessarily sharing information to gain access. These concerns are highly contextual and depend on access needs (Akter et al. 2020; Alharbi, Brewer, and Schoenebeck 2022; Stangl et al. 2020).

The necessary sharing of personal information reflects “forced intimacy,” where “disabled people [are] expected to share personal parts of [them]selves to survive in an ableist world” (Mingus 2017). In this context, users without accessible alternatives must disclose personal information (textual or visual) through Be My Eyes to accomplish tasks and meet access needs.

Our work both supports and extends these privacy considerations through the first analysis of Be My Eyes' reviews, TOS, and privacy policies, revealing additional concerns around platform-level data collection and use beyond those studied in prior VDS research.

(Un)intentional exposure of sensitive information. Be My Eyes' TOS and privacy policy advise users against sharing or displaying anything that may contain private information, including credit cards and mail with personal addresses [TS3, PP4]. This guidance is

problematic for two reasons: (1) given their visual disability, users may not know if such items are in their video feed (Gurari et al. 2019; Akter et al. 2020), and (2) users continue to behave against this advice, potentially because it conflicts with their access needs. Volunteers, for instance, report being asked to read credit card information:

... A woman wanted me to read off a mastercard number she had. I refused as in your privacy policy it states to not provide or give any type of info regarding credit cards ... [P616, Volunteer, 6/15/2021]

Additionally, a user reported using the service to read the last four digits of their debit card and then getting several fraudulent transactions on that card. After reporting, their account was suspended:

... They suspended my account for how long I don't know. They said I violated their policies so now I'm literally in the dark ... [A75, User, 2 yrs ago]

In a TikTok video, the platform's own promotional content features a user using the service with Ray-Ban Meta Smart Glasses to sort mail that contains an address,⁹ despite the platform's policies explicitly advising against such disclosures [TS3, PP4]. Overall, this conflict between users' access needs and platform policies highlights the need for technical interventions.

Use of video content in model training. Be My Eyes' privacy policy broadly permits sharing of the video streams from user-volunteer calls:

We may also provide copies of videos to educational institutions, non-profit organizations, or other companies, including for-profit companies, that are performing research or working to develop products and services, including artificial intelligence systems, that may support blind and low-vision people or other members of the public. [PP4]

A similar passage appears in all four versions of the privacy policy we analyzed and hence has appeared in Be My Eyes' privacy policy since at least May 2020. Users cite this data-sharing provision as a barrier to adoption:

Required to sign over content for AI training. I really want to sign up, but can't with a clear conscience do it when you're forced to sign over your voice for possible use in AI training. Let me opt out if AI training based on my content and I'm in... [P28, 6/16/2025]

Responding to such concerns, Be My Eyes added an opt-out in July 2024:

We provide a way for users to opt out of our announced initiative to provide video data to organizations so they can better train their AI models to reflect

9. <https://www.tiktok.com/@bemyeyesapp/video/7552895781456317718>

the lived experiences of blind people. [PP3]

Users can opt out by emailing the company or navigating to Settings → Terms & Privacy Policy → Video sharing. However, the language of this policy and the multi-step opt-out process create barriers for users who wish to control their data, and an opt-in approach would be more user-centered and consent-focused.

Data sharing with third parties. Volunteers and users have expressed concern about the company's data-sharing practices:

... I am uncomfortable with their privacy policies. It seems like they want to collect a lot of personal information and they even say they will give your information to other organizations ... [A287, Volunteer, 7 yrs ago]

I am blind. This is just creepy. I read the privacy policy. No privacy effectively ... [P773, User, 1/28/2020]

These practices have fluctuated over time across policies, with the most recent one specifying:

Be My Eyes does not sell the personal data we collect to third parties for advertising, marketing, or tracking... Our reasons for using and sharing information are to create and provide better services for people who are blind or have low vision. That includes ... helping other companies create better and more inclusive AI models, tools, and image recognition technologies ... [and] sharing information with research, educational, non-profit, for-profit, or other organizations to give them a better window into providing services and tools for blind and low-vision individuals. [PP4]

While this policy appears to restrict data sharing more than previous iterations, the presence of any data sharing with third parties may raise concern, particularly for disabled people. Marks (2020) observes that wide-scale collection, aggregation, and analysis of people's data can generate algorithmically derived inferences about their disabilities, called "emergent medical data" (EMD), without their consentful disclosure. This information can then enable discrimination by advertisers, employers, or other organizations.

Recent legal cases illustrate these risks. In *CWA v. Facebook* (2018), the platform settled claims that it enabled advertisers to unlawfully target users based on disability and other protected characteristics (*CWA v. Facebook* 2018). The ongoing *Mobley v. Workday* case alleges that the hiring platform's algorithms discriminate based on race, age, and disability (*Mobley v. Workday* 2024). These cases demonstrate how data collected and shared by accessibility platforms can be weaponized against the disabled users they aim to serve.

Although data is not explicitly sold to advertisers, sharing data with other organizations may still enable third parties to infer app usage patterns and, by extension, disability status. Thus, our findings suggest that (1) Be My Eyes is collecting and sharing data that could contribute to EMD and (2) users and volunteers have expressed concerns about these data practices.

Trust and safety challenges. Prior work has explored the problems that users (1) cannot reliably know, and therefore cannot reliably control, what sensitive information is visible to a human- or AI-powered VDS and (2) may need to make such information visible to gain access to it (e.g., Stangl et al. 2020; Alharbi, Brewer, and Schoenebeck 2022; Ahmed et al. 2015). These problems surface concretely in the studied reviews from both the perspectives of users and volunteers. The tension between solutions to these problems raises the challenge of:

Challenge 3: *Accessible privacy protections*

Providing privacy-protecting mechanisms for users to receive visual descriptions of sensitive personal information, recognizing that users may not know what they are holding in front of their cameras and thus cannot precisely control what sensitive information they expose.

Reviews also show that this lack of control extends to the platform itself. Video content is shared for AI training by default, and the opt-out process introduced in 2024 requires users to navigate multiple settings or email support. This default-in approach to a sensitive, disability-linked dataset raises the challenge of:

Challenge 4: *Meaningful participation in AI training*

Accounting for the need for diverse, representative training data to improve AI accessibility tools with individual privacy rights and meaningful consent.

Beyond model training, our analysis of the platform's data-sharing practices—and the emergent medical data risks this can create for a population already vulnerable to disability-based discrimination—points to a related, broader challenge:

Challenge 5: *Accounting for users' privacy threats and needs*

Balancing benefits from data sharing with the privacy rights and safety of disabled users who face heightened risks and may have limited alternative options.

5.3 Be My AI

Sections 5.1 and 5.2 discussed challenges with Be My Eyes' volunteer-powered VDS, the platform's longest-running feature. In this section, we examine Be My AI, a feature announced in 2023 that uses generative AI to provide image descriptions, with plans for video descriptions [TS3].

As discussed in Section 5.2, the use of user- and volunteer-generated content to train AI models is a key privacy consideration. Be My Eyes' July 2024 privacy policy update prohibits using inputs to Be My AI for model training [PP3], reversing an earlier policy that permitted such use [PP2]. While this represents a marked improvement in user privacy and autonomy, the protection applies only to inputs to Be My AI; video streams from user-volunteer calls remain available for training purposes by default. With this training data concern addressed for Be My AI, we turn to other challenges related to the service.

Inaccuracies & hallucinations. Through diary studies (Gonzalez Penuela et al. 2024), interviews (Alharbi et al. 2024), and an exploratory study (Chang et al. 2025), prior work has shown that inaccuracies and hallucinations in AI-generated image and video descriptions negatively impact blind users' experiences. We find similar reports surface publicly in app reviews:

... I scanned the menu with my phone and asked my AI to describe the menu. it told me that the pancake of the month was a stack of buttermilk pancakes with maple syrup and a dollop of butter. But when the waiter brought up my food, she brought out a stack of pancakes with raspberry sauce on top ... [A32, User, 1 yr ago]

Hallucinations are inherent to current large language models (Xu, Jain, and Kankanhalli 2025; Alansari and Luqman 2026; Kalai et al. 2025), raising questions about the appropriateness of AI-generated visual descriptions for certain tasks where visual description accuracy is needed, such as reading allergens listed on print menus. Wang, Kameswaran, and Kacorri (2025) explain that this creates disproportionate harm for BLV people, who cannot visually verify whether AI-generated descriptions are accurate.

The section of Be My Eyes' TOS focused on Be My AI acknowledges that the system "may sometimes result in incorrect responses that do not accurately describe real people, places, or facts" and, consequently, explicitly disallows users from using the AI service for critical issues "like reading a prescription dosage" [TS3]; however, similar to what was observed in Section 5.2 regarding the privacy advice for the volunteer-powered service, users violate this policy to meet access needs:

... I used it this morning to read a prescription label. It did not read the prescription that was in front of me. It seems to be holding onto memories of reading a previous prescription. It's very dangerous ... [A3, User, <1 yr ago]

Conflicting guardrails & user needs. Users express frustration with content guardrails that limit what the AI system will describe. For example:

... I have a major problem with the amount of censorship it has. Everything that we would be trying to get described to us, sighted people would be noticing...

[A37, User, 1 yr ago]

As A37 notes, these limitations can prevent users from accessing the very information they need the service to provide. This echoes Alharbi, Brewer, and Schoenebeck's (2022) findings that nonconsensual information obfuscation is experienced as a form of oppression and ableism.

Some of these limitations relate to identity recognition. For example, the system may not identify gender:

... it is very annoying because it will not tell you if you have a picture of a man or a woman ... even if you're taking a selfie And you know that you're a man, in your head and in real life. [A12, User, <1 yr ago]

This restriction highlights the tension between access needs and concerns about how automated recognition systems misinterpret and stereotype identities. In the case of gender, automated recognition systems operationalize gender in binary, trans-exclusive ways (Keyes 2018). Prior work demonstrates that people want contextual control over how their identities are described and problematizes how automated identity descriptions may provide needed access for some users while perpetuating harm toward other marginalized groups through biased, inaccurate descriptions (Bennett et al. 2021).

These issues are compounded by inconsistent guardrail performance:

... I use the app mainly for reading mail. Occasionally, it will tell me that it can't read the recipients name or address, because it is not supposed to disclose information from an image. To me, that kind of defeats the purpose of using the app for reading mail as a blind person. Yet, other times, it will read the mail, including giving me the name and recipients address. So the inconsistent application of this rule is confusing. But ultimately, the real frustration is that it applies this rule in the first place, meaning that I can't always use this app, because I have to switch to another app to read the image... [A29, User, 1 yr ago]

The unreliable application of these guardrails creates additional barriers that users must work around to gain access to information.

Facial processing violating bystander (and user) privacy. To comply with biometric privacy laws, Be My Eyes' TOS requires users to obtain consent before knowingly submitting images containing faces:

If you intentionally submit an image of other people to our AI services, you represent and warrant to us that each such person has, in compliance with any applicable laws, given their permission and consent to you, to Be My Eyes, and to our third-party AI provider to process that image, including in ways that

may be considered biometric scanning in any applicable jurisdiction...

If you do not consent to the processing of your face as described above, or if you do not have the permission of other people to process images of their face(s) as described above, then you agree not to upload any images of those faces, including your own. [TS3]

However, users are often unable to determine whether faces—including their own—appear in their images, meaning most facial submissions may be unintentional and thus exempt from the consent requirement. Furthermore, integrations with wearable cameras like Ray-Ban Meta Smart Glasses increase exposure risk, as bystanders may not realize they are in frame. While this policy may protect the company from liability, it provides no practical mechanism for users or bystanders to meaningfully consent to or decline facial processing, violating their privacy (Denning, Dehlawi, and Kohno 2014).

On the other hand, privacy compliance for facial processing can conflict with user access needs, such as identifying people for social or safety reasons (Ahmed et al. 2015; Ahmed et al. 2016). This became apparent during 2023 beta testing, when Be My Eyes acknowledged:

Be My AI (powered by OpenAI's GPT-4 image-to-text model) has proven to be an effective tool for users who are blind or have low vision. However, as part of beta testing, OpenAI's engine may blur faces or omit facial descriptions if a person is in the picture because of legitimate privacy considerations. The resulting lack of facial description is frustrating and disappointing to blind and low vision users who hope that AI will provide the same helpful information a sighted person obtains from just looking around. We share that sentiment. (Be My Eyes 2023a)

These competing concerns between bystander privacy and user access needs create a fundamental design challenge for VDS platforms.

Trust and safety challenges. Building on the challenge of providing users with access to information in a privacy-preserving manner (Challenge 3), such access must also be trustworthy. Reported errors from Be My AI thus raise the following:

Challenge 6: *Accounting for hallucinations*

Addressing the heightened risks of AI hallucinations for users, who face additional barriers in validating the accuracy of AI-generated descriptions, particularly for tasks requiring accuracy.

Beyond factual inaccuracies, we observe how content guardrails can conflict with users' access needs, highlighting the challenge of:

Challenge 7: *Non-obstructive AI-output moderation*

Determining what AI-generated descriptions should and should not describe, considering user autonomy and access needs as well as safety and ethics.

Facial processing conducted while using Be My AI presents a challenge at the intersection of privacy and AI, especially given discreet wearable integrations like Ray-Ban Meta glasses:

Challenge 8: *Bystander biometric privacy*

Providing users with the visual information they need while recognizing that they may not be able to identify bystanders in their camera view and therefore may not be able to obtain meaningful consent without platform-level assistance.

Taken together, these challenges show that Be My AI does not resolve the trust and safety tensions raised by the volunteer-powered service (Sections 5.1 and 5.2) so much as echo them: hallucinations and guardrails mirror online safety issues as threats to reliability and access, while biometric processing mirrors volunteer-user disclosure as a central, non-trivial privacy issue.

6 Discussion & Recommendations

Building on the eight trust and safety challenges we identify relating to accessibility, security, privacy, and AI (Section 5), we provide guidance and future directions for VDS development. Scholars and activists have called for honoring disabled people's situated knowledge and following their leadership in developing assistive technologies (Williams et al. 2021; Sins Invalid, n.d.; Hamraie and Fritsch 2019), commonly articulated through the principle "nothing about us without us." Our recommendations should be implemented in ways that honor this principle and should be considered not as comprehensive solutions but as pathways toward safer, more trustworthy VDSs that support disabled people's autonomy and access.

6.1 Expanded Threat Models

Prior work has considered threat models related to both intentional and unintentional exposure of sensitive information (e.g., medical information, bystanders' faces, "messy" self-presentation) that BLV people face when using VDSs powered by AI systems, volunteers, or paid workers (Akter et al. 2020; Alharbi, Brewer, and Schoenebeck 2022; Stangl et al. 2020). Other work has considered how VDSs can be used to mitigate threats related to physical security (Ahmed et al. 2015; Ahmed et al. 2016; Branham et al. 2017). Our findings expand these threat models involving VDSs and contribute new ones, which inform future developments and interventions.

Online safety threats. Challenges 1 (robust, non-obstructive safeguards) and 2 (trauma-informed reporting and support) highlight how online safety threat models, often considered in other contexts (Lenhart et al. 2016), are relevant to VDSs. For human-powered VDSs, these threat models consider how *both* the volunteer and user roles could be exploited to cause harm *beyond* exposure of private information. Future work on these challenges should draw on research into these threats from other domains.

For example, in addressing Challenge 1, platforms could implement obscenity detection systems (Liang et al. 2013; Tian et al. 2013; Cheng et al. 2012). However, these systems have variable performance, and implementing them requires careful consideration of how they might introduce additional issues, such as privacy harms from preprocessing technologies or denial of access from false positives. Relevant to Challenge 2, Chen et al. (2022) propose principles for trauma-informed computing—safety, trust, peer support, collaboration, enablement, and intersectionality—which can guide technology design to mitigate rather than create and exacerbate trauma.

Future work should also investigate how user and volunteer identities impact these threats. We did not find clear reports of identity-based discrimination, but online harassment at large has, for example, racialized dynamics that may manifest in this context (Lenhart et al. 2016).

Privacy threats. Challenges 3 (accessible privacy protections) and 8 (bystander biometric privacy) support prior threat models around privacy exposure during interactions; Challenges 4 (meaningful participation in AI training) and 5 (accounting for users' privacy threats and needs) expand these by considering additional ways in which exposure might occur (i.e., data sharing with third parties and the use of data in model training).

For issues around bystander facial processing (Challenge 8), research projects have developed systems that crop bystanders out of frame using on-device processing on wearable cameras (Corbett et al. 2023) and designed sensors to better indicate to bystanders if they might be in-frame of a wearable (Ahmad et al. 2020). However, these systems are specifically for wearable devices and not mobile phones.

A brief by the Center for Democracy & Technology and the American Association of People with Disabilities outlines “do’s and don’ts” for assistive technology developers (Mishra and Aboulafia 2025), which might provide guidance in navigating these privacy-related threat models. For example, their guidance “don’t use user data for any other purpose, share it with third parties, or sell it” is currently not met, resulting in Challenges 5 and 4. They also provide the guidance “do incorporate a transparent and accessible informed notice and consent process,” and we explore how adopting this in response to the privacy threat of Challenge 4 would look in the context of VDSs in Section 6.2.

Disambiguating the spectrum of “AI” harms. The “AI”-related harms documented in our findings, which include hallucination (Challenge 6), guardrails on “inappropriate” ma-

terial (Challenge 7), and the processing of facial data (Challenge 8), are not harms specific to Be My Eyes but rather are broadly relevant to modern AI systems (Wang, Kameswaran, and Kacorri 2025; Crawford 2021). Data collection practices for advertising and the automatic use of user video streams for model training (Section 5.2) constitute additional “AI”-related harms, even when they do not directly involve model inference.

The spectrum of “AI”-related harms above reflects how “AI” is not a discrete technical object but a collection of data extraction infrastructures, various types of algorithmic and computational systems, deployment practices, and business models (Crawford 2021), hence our use of quotation marks around “AI” in this discussion section. Thus, threat modeling and development of “Responsible AI” in the context of VDSs must contend with the numerous ways data collection, model training and system deployments may harm users and volunteers.

Future work. Future investigations should continue to expand threat models for VDSs to consider risks that were not captured by our analysis. For example, volunteers’ social media posts about positive experiences on the app—which can be an effective way to grow the volunteer network in a grassroots manner, as evidenced by a TikTok post credited for recruiting 100,000 volunteers (Dunn 2020)—may violate users’ privacy, especially when sensitive interactions are discussed publicly. A 2024 article covers a TikTok video in which a Be My Eyes volunteer describes helping a user examine their partner’s phone to determine if the partner was cheating (Fernandez 2024). In this highly sensitive situation, multiple privacy breaches occurred: the volunteer viewed intimate messages without the partner’s knowledge or consent, and then shared details of this private situation with a public online audience. While this content likely was not intended to cause harm, it demonstrates the need for improved protection for sensitive interactions and information.

Additionally, future work should explore related visual description services, such as Aira (Aira, n.d.), which provides assistance from paid workers rather than volunteers. We did not include Aira or other services because our study focused on Be My Eyes, given its relative size, range of features, and relevance as a VDS. Comparing user experiences and policies across services, as well as investigating how threat models align or vary between them, remains an important direction for future work.

6.2 Developing AI VDSs

Given Challenge 4 (meaningful participation in AI training), the collection and use of data for training AI systems for visual description must be critically examined to ensure that the development of these systems, which may serve BLV people’s access needs, does not compromise people’s privacy or autonomy.

Critical refusal is the practice of investigating power imbalances and actively, deliberately creating and moving toward more just alternatives (Garcia et al. 2022). *The Feminist*

Data Manifest-No (Cifor et al. 2019) is a set of 32 commitments that employ critical refusal to reimagine the status quo of data practices. Garcia et al. (2022) explore how these commitments can be realized in practice, providing guidance for reimagining Be My Eyes' data practices in response to Challenge 5 (accounting for users' privacy threats and needs) and Challenge 4. For example, commitment 15 states:

We refuse systems that simplify consent into a one-time action, a simple click of a yes to a terms of service agreement, to ownership of our data in perpetuity. We commit to enacting Planned Parenthood's FRIES model of consent that ensures that it is always "Freely given, Reversible, Informed, Enthusiastic and Specific." (Cifor et al. 2019)

Applying this to Be My Eyes reveals how consent practices could be transformed and, consequently, future directions for research and development. Rather than one-time opt-out mechanisms for video content used in model training (Section 5.2), the FRIES model would enable ongoing, granular consent that users could freely grant or revoke. This could look like users and volunteers deciding on a per-call basis whether their video may be used for training. Future work should investigate how to conceptualize and operationalize the informed aspect of FRIES in this context, given a user might not know the entirety of what is in a video or image. Additionally, requiring FRIES-compliant consent for data sharing would allow users who may have concerns about emergent medical data and algorithmic bias to protect their data from third parties.

Given Be My Eyes' commitment to training models that serve BLV users (Section 2.1), Theodorou et al. (2021) present guidance for "disability-first dataset creation" based on their experiences constructing an object recognition dataset with BLV data collectors. They discuss challenges related to including disabled communities while respecting the invisible labor constantly and broadly demanded of them in daily life. Future work toward this goal should continue to reflexively engage with these challenges and the commitments of *The Feminist Data Manifest-No* (Cifor et al. 2019) to build accessible technologies that respect the data autonomy and dignity of disabled communities.

6.3 Human-Powered Alternatives to AI

Given the AI-related challenges and concerns, as well as other considerations like AI service costs (Chang et al. 2025), environmental impacts (Strubell, Ganesh, and McCallum 2019; Bender et al. 2021), and problematic identity descriptions (Keyes 2018; Bennett et al. 2021), some users may prefer solutions that do not use AI. We argue that Be My Eyes has untapped volunteer capacity to support human-powered alternatives, as volunteers consistently report wanting more opportunities to assist.

Among reviews, reports of insufficient call volume and infrequent matching with users were the second most common issue after technical problems. For example:

This app is such a cool idea and fills a much needed gap. I just wish it would

limit the amount of volunteers so there was a more equal ratio of volunteers to people needing help. Or perhaps introduce another way for all these volunteers to help so that we are being utilized more! [A124, Volunteer, 4yrs ago]

Alternate modalities would also allow more volunteers to participate. For example, a Deaf volunteer stated:

I wanna help them but I am Deaf, so I can't speak to them how can I help them [P41, (Potential) Volunteer, 5/28/2025]

This overwhelming demand for volunteer opportunities and interest in new modalities positions Be My Eyes to implement large-scale, crowd-powered solutions. One direction is CrowdMask, “a system that allows images with potentially sensitive content to be masked by appearing in progressively larger, more identifiable segments, with volunteers masking portions as soon as risks are identified” (Kaur et al. 2017), which to our knowledge has not yet been applied to real-world VDSs. Adapting this for Be My Eyes would allow users to receive privacy-preserving image descriptions through volunteers who identify sensitive details in small photo segments or describe masked images.

We do not propose that such human-powered systems will evade all trust and safety challenges. However, they may offer a way to better meet the needs and desires of users in ways AI-powered ones cannot.

7 Conclusion

We examine trust and safety challenges at the intersection of accessibility, security, privacy, and AI by analyzing 1,428 one-, two-, and three-star reviews of the Be My Eyes app, along with the platform’s terms of service and privacy policies. We chose to focus on a single app, Be My Eyes, based on its relative size, range of features, and relevance as a VDS. Although app-specific features, policies, and user demographics may not reflect all VDSs, we argue that the core trust and safety challenges we identify manifest across VDS platforms, even if the specific implementations differ. We identify issues related to online harassment and abuse (e.g., trolling, fraud, and sexual harassment), as well as weaknesses in reporting systems. We also highlight privacy risks associated with accessing and potentially exposing sensitive information, including through data sharing with third parties and the use of data for AI model training. In addition, we surface issues specific to Be My AI, such as hallucinations, tensions between guardrails and user needs, and concerns related to bystander privacy and facial data. Based on these challenges, we offer recommendations for future development and research, including expanded threat models that reflect our findings, critical examinations of how AI-powered VDSs are designed and deployed, and exploration of human-powered alternatives to AI systems. Through this work, we aim to strengthen trust and safety across VDSs.

References

- Administration for Community Living. 2025. “Apps for People Who Are Blind or Who Have Low Vision,” July 23, 2025. <https://acl.gov/ada/apps-people-who-are-blind-or-who-have-low-vision>.
- Ahmad, Imtiaz, Rosta Farzan, Apu Kapadia, and Adam J. Lee. 2020. “Tangible Privacy: Towards User-Centric Sensor Designs for Bystander Privacy.” *Proceedings of the ACM on Human-Computer Interaction* 4 (CSCW2). <https://doi.org/10.1145/3415187>.
- Ahmed, Tousif, Roberto Hoyle, Kay Connelly, David Crandall, and Apu Kapadia. 2015. “Privacy Concerns and Behaviors of People with Visual Impairments.” In *Proceedings of the 33rd Annual ACM Conference on Human Factors in Computing Systems*, 3523–32. <https://doi.org/10.1145/2702123.2702334>.
- Ahmed, Tousif, Patrick Shaffer, Kay Connelly, David Crandall, and Apu Kapadia. 2016. “Addressing Physical Safety, Security, and Privacy for People with Visual Impairments.” In *Twelfth Symposium on Usable Privacy and Security (SOUPS 2016)*, 341–54. <https://www.usenix.org/conference/soups2016/technical-sessions/presentation/ahmed>.
- Aira. n.d. “Aira.” <https://aira.io>.
- Akter, Taslima, Bryan Dosono, Tousif Ahmed, Apu Kapadia, and Bryan Semaan. 2020. “‘I Am Uncomfortable Sharing What I Can’t See’: Privacy Concerns of the Visually Impaired with Camera Based Assistive Applications.” In *29th USENIX Security Symposium (USENIX Security 20)*, 1929–48. <https://www.usenix.org/conference/usenixsecurity20/presentation/akter>.
- Alansari, Aisha, and Hamzah Luqman. 2026. “Large Language Models Hallucination: A Comprehensive Survey.” *arXiv*, <https://doi.org/10.48550/arXiv.2510.06265>.
- Alharbi, Rahaf, Robin N. Brewer, and Sarita Schoenebeck. 2022. “Understanding Emerging Obfuscation Technologies in Visual Description Services for Blind and Low Vision People.” *Proceedings of the ACM on Human-Computer Interaction* 6 (CSCW2). <https://doi.org/10.1145/3555570>.
- Alharbi, Rahaf, Angela D. Cheong, Jaylin Herskovitz, Robin N. Brewer, and Sarita Schoenebeck. 2025. “‘Trying to Piece It Together’: Exploring Accessible Error Detection in Emerging Privacy Techniques with Blind People.” In *Proceedings of the 27th International ACM SIGACCESS Conference on Computers and Accessibility*. <https://doi.org/10.1145/3663547.3746376>.
- Alharbi, Rahaf, Pa Lor, Jaylin Herskovitz, Sarita Schoenebeck, and Robin N. Brewer. 2024. “Misfitting with AI: How Blind People Verify and Contest AI Errors.” In *Proceedings of the 26th International ACM SIGACCESS Conference on Computers and Accessibility*. <https://doi.org/10.1145/3663548.3675659>.

- Avila, Mauro, Katrin Wolf, Anke Brock, and Niels Henze. 2016. "Remote Assistance for Blind Users in Daily Life: A Survey about Be My Eyes." In *Proceedings of the 9th ACM International Conference on Pervasive Technologies Related to Assistive Environments*. <https://doi.org/10.1145/2910674.2935839>.
- Bardzell, Shaowen, and Jeffrey Bardzell. 2011. "Towards a Feminist HCI Methodology: Social Science, Feminism, and HCI." In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*, 675–84. <https://doi.org/10.1145/1978942.1979041>.
- Be My Eyes. 2021a. "Be My Eyes Wins an Apple Design Award," June 16, 2021. <https://www.bemyeyes.com/news/be-my-eyes-wins-an-apple-design-award/>.
- . 2021b. "Six Years of Be My Eyes," January 20, 2021. <https://www.bemyeyes.com/news/six-years-of-be-my-eyes>.
- . 2023a. "Image to Text AI and Blurred Faces: What's Going On? An Important Update from Be My Eyes about Our New Image to Text AI Features and Facial Description," September 18, 2023. <https://www.bemyeyes.com/blog/image-to-text-ai-and-blurred-faces-whats-going-on/>.
- . 2023b. "Introducing Be My AI (Formerly Virtual Volunteer) for People Who Are Blind or Have Low Vision, Powered by OpenAI's GPT-4," March 14, 2023. <https://www.bemyeyes.com/news/introducing-be-my-ai-formerly-virtual-volunteer-for-people-who-are-blind-or-have-low-vision-powered-by-openais-gpt-4/>.
- . 2024a. "Be My Eyes and Meta Announce Accessibility Partnership," September 25, 2024. <https://www.bemyeyes.com/news/be-my-eyes-and-meta-announce-accessibility-partnership>.
- . 2024b. "Be My Eyes Announces New Collaboration with Microsoft to Help Train Inclusive AI Models," October 17, 2024. <https://www.bemyeyes.com/news/be-my-eyes-announces-new-collaboration-with-microsoft-to-help-train-inclusive-ai-models>.
- . 2025. "Be My Eyes Wins Apple's 2025 App Store Award for Cultural Impact," December 4, 2025. <https://www.bemyeyes.com/blog/be-my-eyes-wins-apples-2025-app-store-award-for-cultural-impact/>.
- . n.d.-a. "About us." <https://www.bemyeyes.com/about-us>.
- . n.d.-b. "Be My Eyes." <https://www.bemyeyes.com>.
- Bender, Emily M., Timnit Gebru, Angelina McMillan-Major, and Shmargaret Shmitchell. 2021. "On the Dangers of Stochastic Parrots: Can Language Models Be Too Big?" In *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency*, 610–23. <https://doi.org/10.1145/3442188.3445922>.

- Bennett, Cynthia L., Cole Gleason, Morgan Klaus Scheuerman, Jeffrey P. Bigham, Anhong Guo, and Alexandra To. 2021. “‘It’s Complicated’: Negotiating Accessibility and (Mis)Representation in Image Descriptions of Race, Gender, and Disability.” In *Proceedings of the 2021 CHI Conference on Human Factors in Computing Systems*. <https://doi.org/10.1145/3411764.3445498>.
- Bigham, Jeffrey P., Chandrika Jayant, Hanjie Ji, Greg Little, Andrew Miller, Robert C. Miller, Robin Miller, et al. 2010. “VizWiz: Nearly Real-Time Answers to Visual Questions.” In *Proceedings of the 23rd Annual ACM Symposium on User Interface Software and Technology*, 333–42. <https://doi.org/10.1145/1866029.1866080>.
- Brady, Erin, Meredith Ringel Morris, Yu Zhong, Samuel White, and Jeffrey P. Bigham. 2013. “Visual Challenges in the Everyday Lives of Blind People.” In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*, 2117–26. <https://doi.org/10.1145/2470654.2481291>.
- Branham, Stacy M., Ali Abdolrahmani, William Easley, Morgan Scheuerman, Erick Ronquillo, and Amy Hurst. 2017. “‘Is Someone There? Do They Have a Gun’: How Visual Information about Others Can Improve Personal Safety Management for Blind Individuals.” In *Proceedings of the 19th International ACM SIGACCESS Conference on Computers and Accessibility*, 260–69. <https://doi.org/10.1145/3132525.3132534>.
- Brief for EEOC as Amicus Curiae, Mobley v. Workday, Inc., No. 3:23-cv-00770-RFL, N.D. Cal., filed April 9, 2024 (2024), <https://www.eeoc.gov/sites/default/files/2024-04/Mobley%20v%20Workday%20NDCal%20am-brf%2004-24%20sjw.pdf>.
- Buck, Amber M., and Devon F. Ralston. 2021. “I Didn’t Sign Up for Your Research Study: The Ethics of Using ‘Public’ Data.” *Computers and Composition* 61. <https://doi.org/10.1016/j.compcom.2021.102655>.
- Burton, Michele A., Erin Brady, Robin Brewer, Callie Neylan, Jeffrey P. Bigham, and Amy Hurst. 2012. “Crowdsourcing Subjective Fashion Advice Using VizWiz: Challenges and Opportunities.” In *Proceedings of the 14th International ACM SIGACCESS Conference on Computers and Accessibility*, 135–42. <https://doi.org/10.1145/2384916.2384941>.
- Chang, Ruei-Che, Yuxuan Liu, and Anhong Guo. 2024. “WorldScribe: Towards Context-Aware Live Visual Descriptions.” In *Proceedings of the 37th Annual ACM Symposium on User Interface Software and Technology*. <https://doi.org/10.1145/3654777.3676375>.
- Chang, Ruei-Che, Rosiana Natalie, Wenqian Xu, Jovan Zheng Feng Yap, and Anhong Guo. 2025. “Probing the Gaps in ChatGPT’s Live Video Chat for Real-World Assistance for People Who Are Blind or Visually Impaired.” In *Proceedings of the 27th International ACM SIGACCESS Conference on Computers and Accessibility*. <https://doi.org/10.1145/3663547.3746319>.

- Chang, Ruei-Che, Rosiana Natalie, Wenqian Xu, Jovan Zheng Feng Yap, Tiange Luo, Venkatesh Potluri, and Anhong Guo. 2026. "TouchScribe: Augmenting Non-Visual Hand-Object Interactions with Automated Live Visual Descriptions." In *Proceedings of the 2026 CHI Conference on Human Factors in Computing Systems*. <https://doi.org/10.1145/3772318.3791308>.
- Charge of Discrimination, Communications Workers of America v. Facebook, U.S. Equal Employment Opportunity Commission (2018), https://assets.aclu.org/live/uploads/legal-documents/Facebook_bill.pdf.
- Chen, Janet X., Allison McDonald, Yixin Zou, Emily Tseng, Kevin A. Roundy, Acar Tamer-soy, Florian Schaub, Thomas Ristenpart, and Nicola Dell. 2022. "Trauma-Informed Computing: Towards Safer Technology Experiences for All." In *Proceedings of the 2022 CHI Conference on Human Factors in Computing Systems*. <https://doi.org/10.1145/3491102.3517475>.
- Cheng, Hanqiang, Yu-Li Liang, Xinyu Xing, Xue Liu, Richard Han, Qin Lv, and Shivakant Mishra. 2012. "Efficient Misbehaving User Detection in Online Video Chat Services." In *Proceedings of the Fifth ACM International Conference on Web Search and Data Mining*, 23–32. <https://doi.org/10.1145/2124295.2124301>.
- Cifor, Marika, Patricia Garcia, T. L. Cowan, Jasmine Rault, Tonia Sutherland, Anita Chan, Jennifer Rode, Anna Lauren Hoffmann, Niloufar Salehi, and Lisa Nakamura. 2019. "Feminist Data Manifest-No." <https://www.manifestno.com>.
- Corbett, Matthew, Brendan David-John, Jiacheng Shang, Y. Charlie Hu, and Bo Ji. 2023. "BystandAR: Protecting Bystander Visual Data in Augmented Reality Systems." In *Proceedings of the 21st Annual International Conference on Mobile Systems, Applications and Services*, 370–82. <https://doi.org/10.1145/3581791.3596830>.
- Crawford, Kate. 2021. *Atlas of AI: Power, Politics, and the Planetary Costs of Artificial Intelligence*. New Haven: Yale University Press.
- Darwin Holmes, Andrew Gary. 2020. "Researcher Positionality - A Consideration of Its Influence and Place in Qualitative Research - A New Researcher Guide." *Shanlax International Journal of Education* 8 (4): 1–10. <https://doi.org/10.34293/education.v8i4.3232>.
- Denning, Tamara, Zakariya Dehlawi, and Tadayoshi Kohno. 2014. "In Situ with Bystanders of Augmented Reality Glasses: Perspectives on Recording and Privacy-Mediating Technologies." In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*, 2377–86. <https://doi.org/10.1145/2556288.2557352>.
- Duggan, Maeve. 2017. *Online Harassment 2017*. Pew Research Center. <https://www.pewresearch.org/internet/2017/07/11/online-harassment-2017/>.

- Dunn, Lisa. 2020. "How One Woman Used TikTok to Get 100,000 New Volunteers for an App That Helps Visually Impaired People," November 24, 2020. <https://www.inclusionhub.com/articles/tik-tok-bme>.
- Feal, Álvaro, Paolo Calciati, Narseo Vallina-Rodriguez, Carmela Troncoso, and Alessandra Gorla. 2020. "Angel or Devil? A Privacy Study of Mobile Parental Control Apps." *Proceedings on Privacy Enhancing Technologies* 2020 (2): 314–35. <https://doi.org/10.2478/popets-2020-0029>.
- Fernandez, Stacy. 2024. "'She Hung Up': Woman Uses 'Be My Eyes' App to See If Her Boyfriend Is Cheating on Her," May 29, 2024. <https://dailymail.com/be-my-eyes-app-cheating>.
- Garcia, Patricia, Tonia Sutherland, Niloufar Salehi, Marika Cifor, and Anubha Singh. 2022. "No! Re-imagining Data Practices Through the Lens of Critical Refusal." *Proceedings of the ACM on Human-Computer Interaction* 6 (CSCW2). <https://doi.org/10.1145/3557997>.
- Garland-Thomson, Rosemarie. 2002. "Integrating Disability, Transforming Feminist Theory." *NWSA Journal* 14 (3): 1–32. <https://doi.org/10.2979/nws.2002.14.3.1>.
- Gonzalez Penuela, Ricardo E., Jazmin Collins, Cynthia Bennett, and Shiri Azenkot. 2024. "Investigating Use Cases of AI-Powered Scene Description Applications for Blind and Low Vision People." In *Proceedings of the 2024 CHI Conference on Human Factors in Computing Systems*. <https://doi.org/10.1145/3613904.3642211>.
- Gonzalez Penuela, Ricardo E., Crescentia Jung, Sharon Lin, Ruiying Hu, and Shiri Azenkot. 2026. "How Multimodal Large Language Models Support Access to Visual Information: A Diary Study with Blind and Low Vision People." In *Proceedings of the 2026 CHI Conference on Human Factors in Computing Systems*. <https://doi.org/10.1145/3772318.3793266>.
- Gottfried, Jeffrey, Eugenie Park, and Monica Anderson. 2025. *Online Scams and Attacks in America Today*. Pew Research Center. <https://www.pewresearch.org/internet/2025/07/31/online-scams-and-attacks-in-america-today>.
- Gurari, Danna, Qing Li, Chi Lin, Yinan Zhao, Anhong Guo, Abigale Stangl, and Jeffrey P. Bigham. 2019. "VizWiz-Priv: A Dataset for Recognizing the Presence and Purpose of Private Visual Information in Images Taken by Blind People." In *2019 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 939–48. <https://doi.org/10.1109/CVPR.2019.00103>.
- Gurari, Danna, Qing Li, Abigale J. Stangl, Anhong Guo, Chi Lin, Kristen Grauman, Jiebo Luo, and Jeffrey P. Bigham. 2018. "VizWiz Grand Challenge: Answering Visual Questions from Blind People." In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 3608–17. <https://doi.org/10.1109/CVPR.2018.00380>.

- Hamraie, Aimi, and Kelly Fritsch. 2019. "Crip Technoscience Manifesto." *Catalyst: Feminism, Theory, Technoscience* 5 (1). <https://doi.org/10.28968/cftt.v5i1.29607>.
- Hanson, Vicki L., Anna Cavender, and Shari Trewin. 2015. "Writing about Accessibility." *Interactions* 22 (6): 62–65. <https://doi.org/10.1145/2828432>.
- Houtti, Mohamed, Abhishek Roy, Venkata Narsi Reddy Gangula, and Ashley Walker. 2024. "A Survey of Scam Exposure, Victimization, Types, Vectors, and Reporting in 12 Countries." *Journal of Online Trust and Safety* 2 (4). <https://doi.org/10.54501/jots.v2i4.204>.
- Kalai, Adam Tauman, Ofir Nachum, Santosh S. Vempala, and Edwin Zhang. 2025. "Why Language Models Hallucinate." *arXiv*, <https://doi.org/10.48550/arXiv.2509.04664>.
- Kass, Mirelle, Lindsay M. Alexander, Tatum Connell, Sara Bagheri Hamaneh, Najuma Abdullah, Michelle Freund, Michael P. Milham, and Aki Nikolaidis. 2026. "Negative Online Experiences and Reporting Rates in Youth with Mental Health Conditions." *JAACAP Open* 4 (4): 701–12. <https://doi.org/10.1016/j.jaacop.2026.04.001>.
- Kaur, Harmanpreet, Mitchell Gordon, Yiwei Yang, Jeffrey P. Bigham, Jaime Teevan, Ece Kamar, and Walter Lasecki. 2017. "CrowdMask: Using Crowds to Preserve Privacy in Crowd-Powered Systems via Progressive Filtering." *Proceedings of the AAAI Conference on Human Computation and Crowdsourcing* 5 (1): 89–98. <https://doi.org/10.1609/hcomp.v5i1.13314>.
- Keyes, Os. 2018. "The Misgendering Machines: Trans/HCI Implications of Automatic Gender Recognition." *Proceedings of the ACM on Human-Computer Interaction* 2 (CSCW). <https://doi.org/10.1145/3274357>.
- Lenhart, Amanda, Michele Ybarra, Kathryn Zickuhr, and Myeshia Price-Feeney. 2016. *Online Harassment, Digital Abuse, and Cyberstalking in America*. Data & Society Research Institute and Center for Innovative Public Health Research. https://www.datasociety.net/pubs/oh/Online_Harassment_2016.pdf.
- Liang, Calvin A., Sean A. Munson, and Julie A. Kientz. 2021. "Embracing Four Tensions in Human-Computer Interaction Research with Marginalized People." *ACM Transactions on Computer-Human Interaction* 28 (2). <https://doi.org/10.1145/3443686>.
- Liang, Yu-Li, Xinyu Xing, Hanqiang Cheng, Jianxun Dang, Sui Huang, Richard Han, Xue Liu, Qin Lv, and Shivakant Mishra. 2013. "SafeVchat: A System for Obscene Content Detection in Online Video Chat Services." *ACM Transactions on Internet Technology* 12 (4). <https://doi.org/10.1145/2499926.2499927>.
- Marks, Mason. 2020. "Algorithmic Disability Discrimination." In *Disability, Health, Law, and Bioethics*, edited by I. Glenn Cohen, Carmel Shachar, Anita Silvers, and Michael Ashley Stein, 242–54. Cambridge University Press. <https://doi.org/10.1017/9781108622851.026>.

- McGlynn, Clare, and Kelly Johnson. 2021. *Cyberflashing: Recognising Harms, Reforming Laws*. Bristol University Press. <https://doi.org/10.2307/j.ctv1jk0hd4>.
- Microsoft. n.d.-a. "Seeing AI." <https://www.seeingai.com>.
- . n.d.-b. "Seeing AI." <https://www.microsoft.com/en-us/garage/wall-of-fame/seeing-ai/>.
- Mingus, Mia. 2017. "Forced Intimacy: An Ableist Norm," August 6, 2017. <https://leavingevidence.wordpress.com/2017/08/06/forced-intimacy-an-ableist-norm>.
- Mishra, Achi, and Ariana Aboulafia. 2025. *Inclusive Innovation: How to Incorporate Privacy into Inclusive Design for Assistive Technologies*. Center for Democracy & Technology. <https://cdt.org/insights/inclusive-innovation-how-to-incorporate-privacy-into-inclusive-design-for-assistive-technologies/>.
- Owens, Kentrell, Anita Alem, Franziska Roesner, and Tadayoshi Kohno. 2022. "Electronic Monitoring Smartphone Apps: An Analysis of Risks from Technical, Human-Centered, and Legal Perspectives." In *31st USENIX Security Symposium (USENIX Security 22)*, 4077–94. <https://www.usenix.org/conference/usenixsecurity22/presentation/owens>.
- Robertson, Ronald E. 2022. "Uncommon Yet Consequential Online Harms." *Journal of Online Trust and Safety* 1 (3). <https://doi.org/10.54501/jots.v1i3.87>.
- Sharif, Ather, Aedan Liam McCall, and Kianna Roces Bolante. 2022. "Should I Say 'Disabled People' or 'People with Disabilities'? Language Preferences of Disabled People Between Identity- and Person-First Language." In *Proceedings of the 24th International ACM SIGACCESS Conference on Computers and Accessibility*. <https://doi.org/10.1145/3517428.3544813>.
- Sharma, Tanusree, Yu-Yun Tseng, Lotus Zhang, Ayae Ide, Kelly Avery Mack, Leah Findlater, Danna Gurari, and Yang Wang. 2025. "'Before, I Asked My Mom, Now I Ask ChatGPT': Visual Privacy Management with Generative AI for Blind and Low-Vision People." In *Proceedings of the 27th International ACM SIGACCESS Conference on Computers and Accessibility*. <https://doi.org/10.1145/3663547.3746335>.
- Sins Invalid. n.d. "10 Principles of Disability Justice." <https://sinsinvalid.org/10-principles-of-disability-justice>.
- Stangl, Abigale, Kristina Shiroma, Bo Xie, Kenneth R. Fleischmann, and Danna Gurari. 2020. "Visual Content Considered Private by People Who Are Blind." In *Proceedings of the 22nd International ACM SIGACCESS Conference on Computers and Accessibility*. <https://doi.org/10.1145/3373625.3417014>.
- Strubell, Emma, Ananya Ganesh, and Andrew McCallum. 2019. "Energy and Policy Considerations for Deep Learning in NLP." In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, 3645–50. <https://doi.org/10.18653/v1/P19-1355>.

- Theodorou, Lida, Daniela Massiceti, Luisa Zintgraf, Simone Stumpf, Cecily Morrison, Edward Cutrell, Matthew Tobias Harris, and Katja Hofmann. 2021. "Disability-first Dataset Creation: Lessons from Constructing a Dataset for Teachable Object Recognition with Blind and Low Vision Data Collectors." In *Proceedings of the 23rd International ACM SIGACCESS Conference on Computers and Accessibility*. <https://doi.org/10.1145/3441852.3471225>.
- Tian, Lei, Shaosong Li, Junho Ahn, David Chu, Richard Han, Qin Lv, and Shivakant Mishra. 2013. "Understanding User Behavior at Scale in a Mobile Video Chat Application." In *Proceedings of the 2013 ACM International Joint Conference on Pervasive and Ubiquitous Computing*, 647–56. <https://doi.org/10.1145/2493432.2493488>.
- Wang, Lining, Vaishnav Kameswaran, and Hernisa Kacorri. 2025. "Toward a Taxonomy of Algorithmic Harms for Disability: A Systematic Review." *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society* 8 (3): 2649–65. <https://doi.org/10.1609/aies.v8i3.36745>.
- Williams, Rua M., Kathryn Ringland, Amelia Gibson, Mahender Mandala, Arne Maibaum, and Tiago Guerreiro. 2021. "Articulations toward a crip HCI." *Interactions* 28 (3): 28–37. <https://doi.org/10.1145/3458453>.
- Xie, Jingyi, Rui Yu, He Zhang, Syed Masum Billah, Sooyeon Lee, and John M. Carroll. 2025. "Beyond Visual Perception: Insights from Smartphone Interaction of Visually Impaired Users with Large Multimodal Models." In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems*. <https://doi.org/10.1145/3706598.3714210>.
- Xu, Ziwei, Sanjay Jain, and Mohan Kankanhalli. 2025. "Hallucination Is Inevitable: An Innate Limitation of Large Language Models." *arXiv*, <https://doi.org/10.48550/arXiv.2401.11817>.
- Zamiri Zeraati, Farnaz, Yang Cao, Yuehan Qiao, Hal Daumé III, and Hernisa Kacorri. 2026. "Say It My Way: Exploring Control in Conversational Visual Question Answering with Blind Users." In *Proceedings of the 2026 CHI Conference on Human Factors in Computing Systems*. <https://doi.org/10.1145/3772318.3791834>.
- Zhang, Lotus, Abigale Stangl, Tanusree Sharma, Yu-Yun Tseng, Inan Xu, Danna Gurari, Yang Wang, and Leah Findlater. 2024. "Designing Accessible Obfuscation Support for Blind Individuals' Visual Privacy Management." In *Proceedings of the 2024 CHI Conference on Human Factors in Computing Systems*. <https://doi.org/10.1145/3613904.3642713>.

Authors

Natalie Grace Brigham is a PhD student at the University of Washington. Email: nbrigham@cs.washington.edu.

Tadayoshi Kohno is a Professor and the Robert L. McDevitt, K.S.G., K.C.H.S. and Catherine H. McDevitt L.C.H.S. Chair in Computer Science at Georgetown University.

Acknowledgments

We thank Lucy Jiang for conducting a sensitivity read, Jennifer Mankoff for their guidance on this topic, and Christina Yeung for her help with qualitative coding. We are also grateful for the thoughtful feedback of Ian Chang, Franzi Roesner, Henry Wong, and our reviewers.

Data and Code Availability

The app reviews analyzed in this study are publicly available on the Apple App Store and Google Play Store. To limit exposure and risk to individual reviewers, we do not share the coded review dataset.

Ethics Statement

Regarding our collection and use of public review data, we received an IRB determination from our institution that this research does not constitute human subjects research under federal regulations. See Section 4 for further discussion of ethical considerations.

Funding Statement

This project was supported in part by NSF grants CNS-2205171 and CNS-2513316 and the University of Washington Tech Policy Lab. Tadayoshi Kohno is supported by the Robert L. McDevitt, K.S.G., K.C.H.S. and Catherine H. McDevitt L.C.H.S. Chair in Computer Science at Georgetown University.

Disclosure Statement

The authors declare no competing interests.

Keywords

Be My Eyes; visual description; blind and low vision users; app reviews; policy analysis